

Pitch video script: Bob Onboarding Pipeline (3:00 cut)

Audience: IBM Bob 2.0 hackathon judges (lablab.ai). Criteria to hit: Application of Technology, Presentation, Business Value, Originality. Speaking pace ~130 words/min. [X] = fill in before recording; do not guess.

Rule check: video $\leq 3:00$, judges stop at 3:00, ≥ 90 s of the solution running. This script targets exactly 3:00, with the live-demo block held at 105 s of screen time — comfortably clears the 90 s floor even though the narration under it is much shorter (see the note after that block).

Source of truth: pulled directly from the speaker notes on the pitch deck (Team Onboard — [deck link](#)), 9 slides. This deck no longer has a separate 5:00 version — the whole thing is now built to 3:00.

What got cut from the earlier 5:00 draft

- The two problem-evidence slides (58%/82% stats, the "docs go stale" diagram) — compressed to one line in the cover beat.
- The "Bob usage" slide (skill/command/MCP-server breakdown) — dropped entirely. If a judge asks about it live, it's in the written 500-word Bob Usage Statement instead.
- The architecture diagram merged into the solution beat instead of its own slide.

3:00 script

Rows 3–7 (`live-intake` through `live-send`) are the live-demo block — 105 s of screen time total, clearing the ≥ 90 s requirement (see the note after the table for why the narration under it looks short).

Time	Slide	On screen	Narration
0:00–0:15	<code>cover</code>	Title card: "Day one, a stale README." Slow / Expensive / Inconsistent	"A new developer joins. They get a stale README, and their senior repeats the same explanations for every hire. Onboarding is slow, expensive, and different every time."
0:15–0:35	<code>solution</code>	<code>/onboard</code> line + flowchart: Intake → Generate ×4 → Verify → Approve → Render → Send	"We built an onboarding pipeline inside Bob IDE. One command, <code>/onboard</code> . Four subagents scan the codebase in parallel, a deterministic check verifies every citation, and nothing renders or sends until a human approves."
0:35–0:55	<code>live-intake</code>	LIVE: Bob asking for the codebase path, then probing it	"Live run. The manager points Bob at a codebase, and Bob confirms the project. Company, brand and repo come from an instance file, so the core stays generic."
0:55–1:20	<code>live-generate</code>	LIVE: four subagents launching in parallel, "10 issues found" badge — hold on this a beat, let the tool calls actually scroll	"Four subagents run in parallel and find ten real issues in this Java app, including hardcoded credentials, each with evidence at file and line. Then a deterministic check verifies every path exists. Anything it can't confirm goes into a flagged list. Bob never invents a fact."
1:20–1:40	<code>live-approval</code>	LIVE: the approval gate, flagged issues + decision options — hold long enough for judges to read the table	"Now the human gate. The manager sees the flagged items and the issue index, and has to explicitly say yes. Auto-approve doesn't count."
1:40–2:00	<code>live-outputs</code>	LIVE: rendered PDF, table of contents and welcome cards	"Approved. Bob renders a branded PDF, a narrated code tour, and a context folder the newcomer loads straight into their own Bob."

Time	Slide	On screen	Narration
2:00–2:20	live-send	LIVE: Ethereal inbox, email received with 3 attachments	"Delivery goes through an email MCP server Bob built itself. The newcomer gets the pack; the manager gets a separate brief."
2:20–2:45	impact	Big numbers: ~20 min vs. 20–40 hrs (LACY, Beko)	"The run took about 20 minutes of machine time and one minute of human approval. Industry data puts manual onboarding at 20 to 40 hours of expert time per hire."
2:45–3:00	close	Closing statement + /build-pipeline · instances/<client> · Team Onboard	"A generic core, a folder per client, a human in control. Thanks."

Why the live block works at only ~65 s of actual talking inside a 105 s slot: the narration is deliberately sparse so judges can watch Bob actually work — the subagents running, the issues table, the PDF opening. Per the recording tips below, record the real run once, then edit: hold each shot for its full window in the table above even where nobody's talking, rather than padding with words. That screen time is what satisfies the ≥90 s rule, not word count.

Recording tips

- Record the live run once, then cut the *waiting* (installs, long tool calls) but keep the *result* on screen for its full window — the table above assumes held shots, not padded narration.
- Show real terminal and Bob UI throughout; judges want proof it runs, not a mock-up.
- If you're short on time on the day: cut from `solution` (skip straight to the diagram, drop the sentence) before you cut anything from the live block — the demo footage is what the ≥90 s rule and the *Application of Technology* criterion actually hinge on.
- `[X]` doesn't appear in this cut — the impact line now uses the sourced 20–40 hrs figure (LACY), not a placeholder.